

ECE ING4 MACHINE LEARNING


Jeremy Cohen



Logistic Regression

Week 2 Review

Machine Learning Process



Classification vs Regression




Regression

What is the temperature going to be tomorrow?

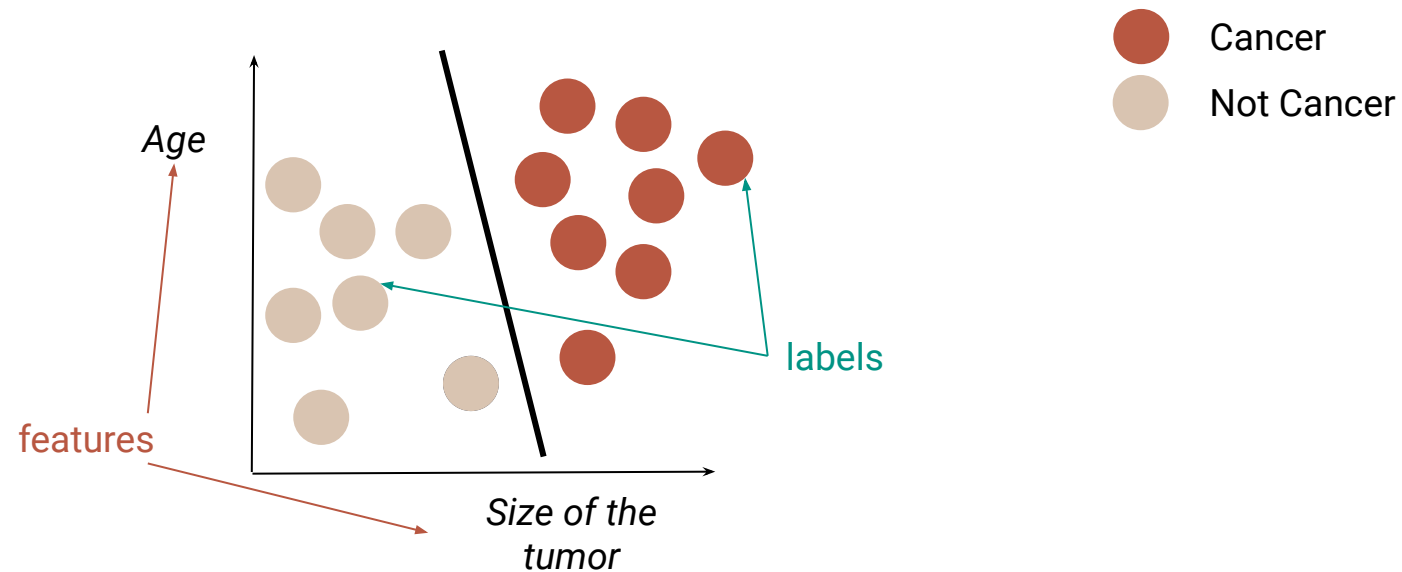


Classification


Will it be Cold or Hot tomorrow?




Classification



Classification



Multi-Class Classification



3 Datasets

Training Set



~70% of the dataset

Used to **train the model**

Validation Set



~20% of the dataset

Used to **test** the model
and **generalize better to
new data**

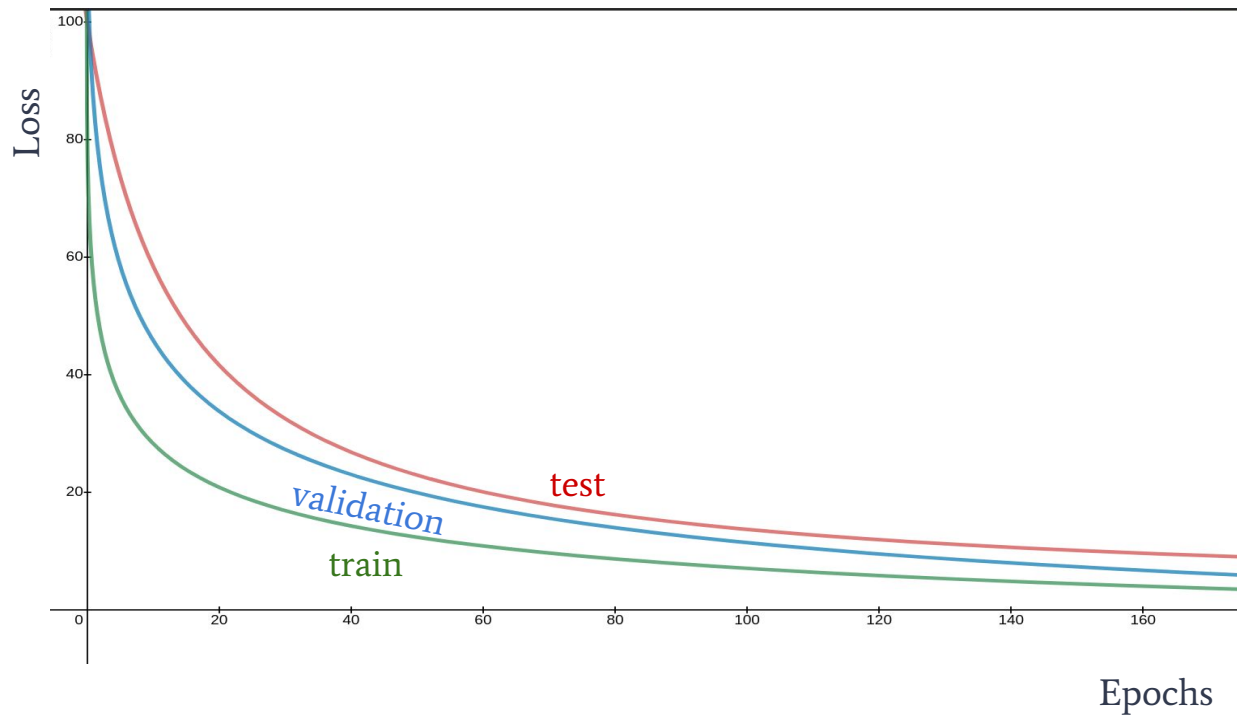
Test Set



~10% of the dataset

Used **only once** when
both accuracies are
good and ready for
real-world

3 Datasets



Linear Regression formula


$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

For convenience of notation, define $x_0 = 1$ ($x_0^{(i)} = 1$)

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$h_{\theta}(x) = \sum_{i=1}^n \theta_i x_i = \theta^T x$$

Mean Squared Error




$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Gradient Descent Recap

$$\theta_1 := \theta_1 - \alpha \frac{\partial}{\partial \theta_1} J(\theta_1)$$

In this case, the derivative
(gradient) $\partial J(\theta) / \partial \theta < 0$.
 $\theta^A - \alpha * \partial J(\theta) / \partial \theta > \theta^A$
 θ^A is moving to the right
 θ is increasing



In this case, the derivative
(gradient) $\partial J(\theta) / \partial \theta > 0$.
 $\theta^B - \alpha * \partial J(\theta) / \partial \theta < \theta^B$
 θ^B is moving to the left
 θ is decreasing

Gradient Descent Recap

Repeat {

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

(simultaneously update θ_j for
 $j = 0, \dots, n$)



}

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_0^{(i)}$$





$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_1^{(i)}$$

$$\theta_2 := \theta_2 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_2^{(i)}$$

...



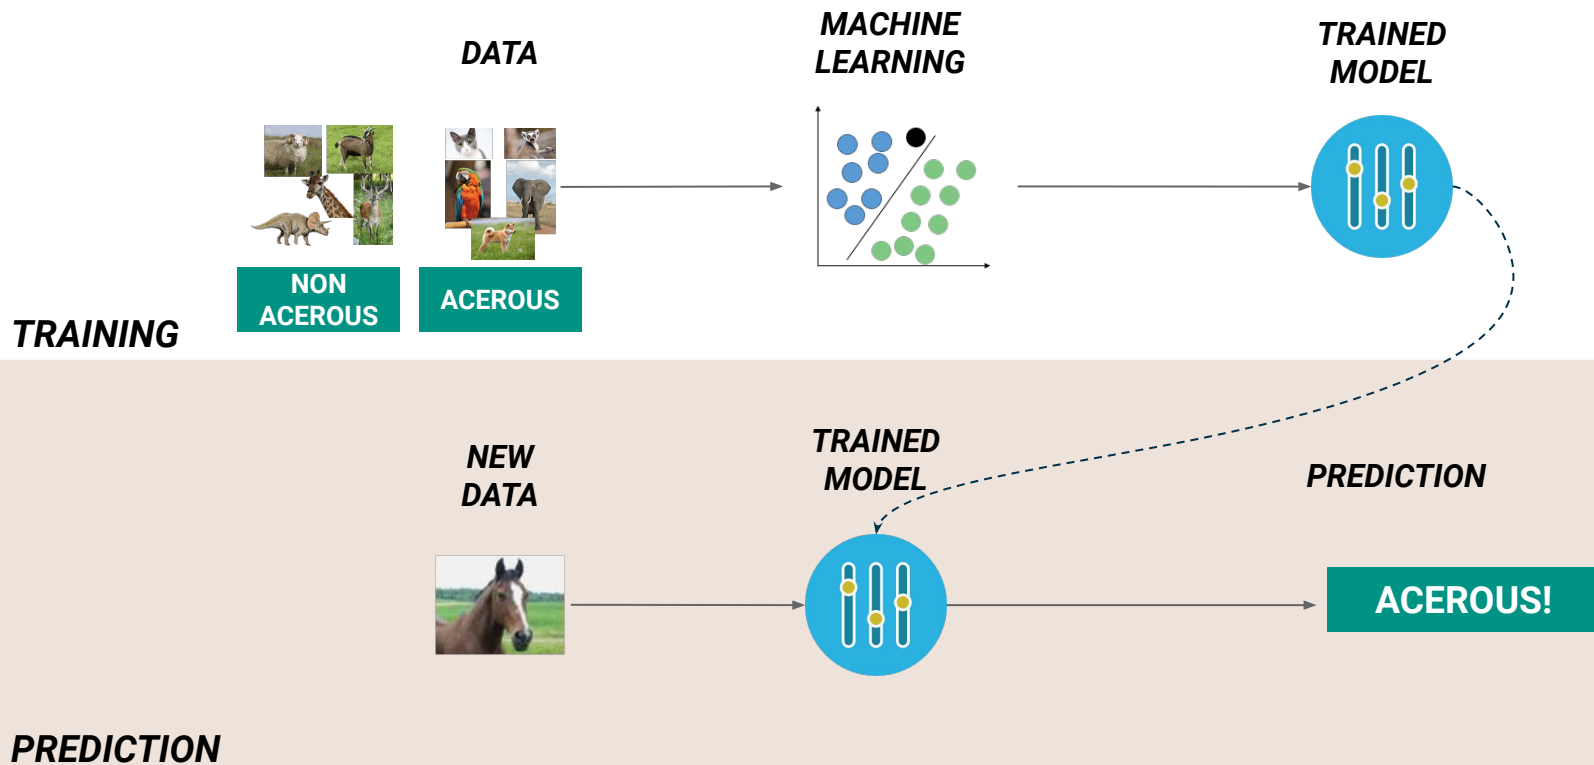
Performance



Classification





Classification Example





Binary vs Multi-Class Classification

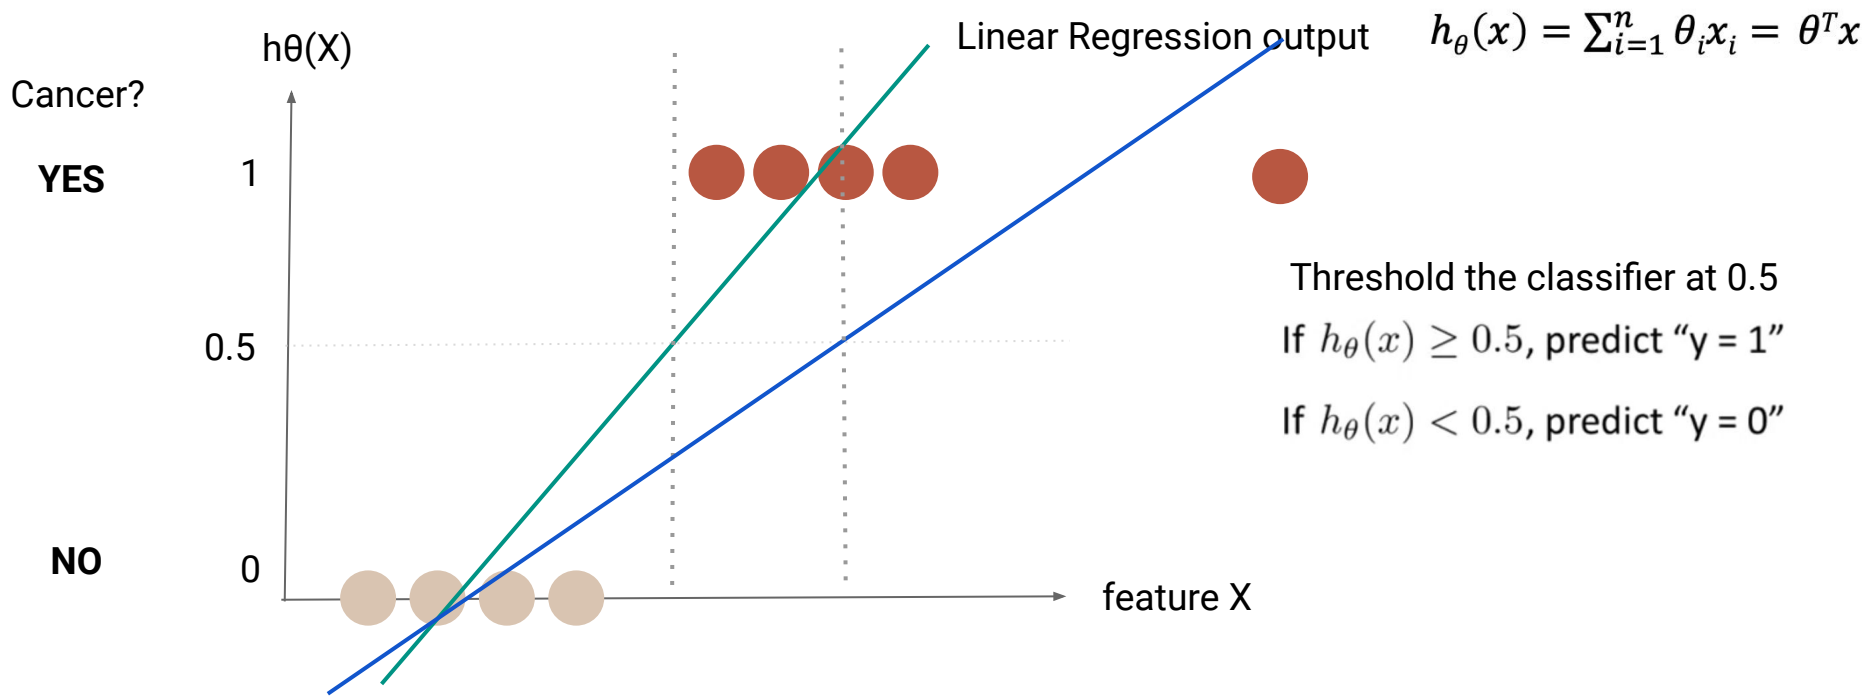
Output has 2 categories



Output has more than categories



Can Linear Regression help?



Linear Regression vs Classification



Linear Regression output : $-\infty < h_{\theta}(x) < +\infty$

Classification output: $y = 0 \text{ or } 1$

Logistic Regression output : $0 \leq h_{\theta}(x) \leq 1$

Linear Regression can't help

We want $0 \leq h_{\theta}(x) \leq 1$



- Use the Sigmoid / Logistic Function : $g(z) = \frac{1}{1 + e^{-z}}$

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

Logistic Regression

$h_{\theta}(x)$ = **estimated probability that $y = 1$** given the input x parameterized by θ


$$h_{\theta}(x) = P(y = 1 | x; \theta) = 1 - P(y = 0 | x; \theta)$$

$$P(y = 1 | x; \theta) + P(y = 0 | x; \theta) = 1$$


Hypothesis:

□ $y = 1$ when $g(\theta^T x) \geq 0.5 \Rightarrow \theta^T x \geq 0$

□ $y = 0$ when $g(\theta^T x) < 0.5 \Rightarrow \theta^T x < 0$



Logistic Regression



$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

If $h_{\theta}(x) \geq 0.5$, predict "y = 1" $\theta^T x \geq 0$

If $h_{\theta}(x) < 0.5$, predict "y = 0" $\theta^T x < 0$


$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2)$$

$$\theta = \begin{bmatrix} -3 \\ 1 \\ 1 \end{bmatrix}$$

Predict "y = 1" if $-3 + x_1 + x_2 \geq 0$

$$x_1 + x_2 \geq 3$$

Non-Linearities



$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

Example: $h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2)$


- $\theta = [-1, 0, 0, 1, 1]$
- Predict 1 if $-1 + x_1^2 + x_2^2 \geq 0 \Rightarrow x_1^2 + x_2^2 \geq 1$
- Predict 0 if $-1 + x_1^2 + x_2^2 < 0 \Rightarrow x_1^2 + x_2^2 < 1$

We can work with non-linearly separable data

Non-Linearities

As with polynomial regression, we can have more complex decision boundaries by adding higher polynomial terms

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2 + \theta_5 x_1^3 + \theta_6 x_2^3)$$




How to choose θ ?

Cost Function

Cost Function for Linear Regression


$$J(\theta_0, \theta_1, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$




$$\frac{1}{1 + e^{-\theta^T x}}$$

Cost Function for Linear Regression

Possible to run Gradient Descent



Impossible to run Gradient Descent



Cost Function for Logistic Regression

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)})$$



$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Cost Function for Logistic Regression

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Cost = 0 if $y = 1, h_{\theta}(x) = 1$
But as $h_{\theta}(x) \rightarrow 0$
 $\text{Cost} \rightarrow \infty$



y=1



Cost Function for Logistic Regression

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

Cost = 0 if $y = 0, h_{\theta}(x) = 0$
But as $h_{\theta}(x) \rightarrow 1$
 $\text{Cost} \rightarrow \infty$



Cost Function for Logistic Regression

$$\text{Cost}(h_{\theta}(x), y) = \begin{cases} -\log(h_{\theta}(x)) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(x)) & \text{if } y = 0 \end{cases}$$

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{Cost}(h_{\theta}(x^{(i)}), y^{(i)}) \\ &= -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

Cost Function for Logistic Regression

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

(simultaneously update all θ_j)

Cost Function for Logistic Regression

Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[\sum_{i=1}^m y^{(i)} \log h_{\theta}(x^{(i)}) + (1 - y^{(i)}) \log (1 - h_{\theta}(x^{(i)})) \right]$$

Want $\min_{\theta} J(\theta)$:

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

(simultaneously update all θ_j)

linear regression


logistic regression

$$h_{\theta}(x) = \theta^T x$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

Support Vector Machine


SVM



GOAL

Separate the dataset into classes with as much correctly classified points as possible


SVM



PROCESS

- Select the line that correctly classifies as many points as possible
- Select the line that maximizes distance with support vectors


SVM



NOTABLE


- Robust to outliers/exceptions; line will not be affected

SVM




NOTABLE


- If the data is non linearly separable, we can make it




Multiclass Classification




MultiClass Classification: One vs All




MultiClass Classification: One vs All



MultiClass Classification: One vs All




MultiClass Classification: One vs All





On a new input, to make a prediction, pick the class that maximizes: $\max_i h_{\theta}^i(x)$

MultiClass Classification: One vs All




On a new input, to make a prediction, pick the class that maximizes: $\max_i h_{\theta}^i(x)$

MultiClass Classification: One vs All




Overfitting

Linear Regression



UNDERFITTING

high bias




JUST RIGHT




OVERFITTING

high variance

Linear Regression




Logistic Regression




UNDERFITTING

high bias



JUST RIGHT



OVERFITTING

high variance


Performance

Accuracy

$$\text{Accuracy} = \frac{\text{number of data points *classified correctly*}}{\text{all data points}}$$

Confusion Matrix



	Positive Examples	Negative Examples
Correct	TRUE POSITIVE	TRUE NEGATIVE
Wrong	FALSE NEGATIVE	FALSE POSITIVE



Imbalanced Classification Problem

If a classifier labels everyone as not terrorist:

Accuracy = 99.9998%




2

x 1,000,000

Precision & Recall

$$\text{Precision} = \frac{\text{True positive}}{\text{True positive} + \text{False positive}}$$

$$\text{Recall} = \frac{\text{True positive}}{\text{True positive} + \text{False negative}}$$



How many selected items are relevant?

$$\text{Precision} = \frac{\text{True positives}}{\text{True positives} + \text{False positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{True positives}}{\text{True positives} + \text{False negatives}}$$

Precision & Recall for a Balanced Dataset

	Positive	Negative
Num examples	53	47
Correct	48	40
Wrong	5	7

$$\text{Accuracy} = \frac{\text{True Positives (48)} + \text{True Negatives (40)}}{\text{All Examples (100)}} = 88\%$$

$$\text{Precision} = \frac{\text{True Positives (48)}}{\text{True Positives (48)} + \text{False Positive (7)}} = 87\%$$

$$\text{Recall} = \frac{\text{True Positives (48)}}{\text{True Positives (48)} + \text{False Negative (5)}} = 90\%$$

Precision & Recall for an imbalanced Dataset




90



10

PREDICTED/ ACTUAL	Positive	Negative
Positive	0 (TP)	0 (FP)
Negative	10 (FN)	90 (TN)



100

$$\text{Accuracy} = \frac{\text{True Positives (0)} + \text{True Negatives (99)}}{\text{All Examples (100)}} = 90\%$$

$$\text{Precision} = \frac{\text{True Positives (0)}}{\text{True Positives (0)} + \text{False Positive (0)}}$$

$$\text{Recall} = \frac{\text{True Positives (0)}}{\text{True Positives (0)} + \text{False Negative (10)}} = 0\%$$

If we label everyone as non-terrorists, our recall and precision are 0 or not computable

Precision & Recall for an imbalanced Dataset




90



10

PREDICTED/ACTUAL	Positive	Negative
Positive	10 (TP)	90 (FP)
Negative	0 (FN)	0 (TN)



100

$$\text{Accuracy} = \frac{\text{True Positives (10)} + \text{True Negatives (0)}}{\text{All Examples (100)}} = 10\%$$

$$\text{Precision} = \frac{\text{True Positives (10)}}{\text{True Positives (10)} + \text{False Positive (90)}} = 10\%$$

$$\text{Recall} = \frac{\text{True Positives (10)}}{\text{True Positives (10)} + \text{False Negative (0)}} = 100\%$$

Recall is the ability to find relevant cases in a dataset!

Precision & Recall for an imbalanced Dataset




90



10

PREDICTED / ACTUAL	Positive	Negative
Positive	3 (TP)	0 (FP)
Negative	7 (FN)	90 (TN)



3



97

$$\text{Accuracy} = \frac{\text{True Positives (3)} + \text{True Negatives (90)}}{\text{All Examples (100)}} = 93\%$$

$$\text{Precision} = \frac{\text{True Positives (3)}}{\text{True Positives (3) + False Positive (0)}} = 100\%$$

Precision is the ability to find only the relevant data points

$$\text{Recall} = \frac{\text{True Positives (3)}}{\text{True Positives (3) + False Negative (7)}} = 30\%$$

Precision & Recall for an imbalanced Dataset




90



10

PREDICTED / ACTUAL	Positive	Negative
Positive	4 (TP)	3 (FP)
Negative	6 (FN)	87 (TN)



13



87

$$\text{Accuracy} = \frac{\text{True Positives (4)} + \text{True Negatives (87)}}{\text{All Examples (100)}} = 91\%$$

$$\text{Precision} = \frac{\text{True Positives (4)}}{\text{True Positives (4)} + \text{False Positive (3)}} = 57\%$$

$$\text{Recall} = \frac{\text{True Positives (4)}}{\text{True Positives (4)} + \text{False Negative (6)}} = 40\%$$

F1 Score

$$F_1 = 2 * \frac{\textit{precision} * \textit{recall}}{\textit{precision} + \textit{recall}}$$


Accuracy is used when the True Positives and True negatives are **more important** while **F1-score** is used when the False Negatives and False Positives are crucial.

Accuracy can be used when the class distribution is similar while **F1-score** is a **better** metric when there are imbalanced classes as in the above case

K-Fold Cross-Validation

Problem

Training Set



Validation Set



A part of the dataset is not used for training

K-Fold Cross Validation

Here, the test set is 20% of the dataset.
We lose $\frac{1}{5}$ of the dataset.



Dataset

Run K-Experiments in which we:

- Choose a test set
- Train
- Test

After all the experiments, we average the results



Thank

You

jeremycohen.podia.com

<https://www.linkedin.com/in/jeremycohen2626/>